

# Multivariate Analysis Final

Richard Noel Britto

12/5/2020

## Multivariate Analysis Final

The accompanied dataset, from Higham et al. (1980), gives 9 skull measurement for different canine groups.

The variables

X1 = length of mandible

X2 = breadth of mandible below 1st molar

X3 = breadth of articular condyle

X4 = height of mandible below first molar

X5 length of first molar, X6 = breadth of first molar

X7 = length of first to third molar inclusive (first to second for Cuon)

X8 = length from first to fourth premolar inclusive

X9 = breadth of lower canine

All measured in millimeters

Please answer the questions below

1. Using suitable graphical method, compare the distribution of the nine variables for the prehistoric and modern Thai dog.

a. Create a Draftsman plot for the 9 variables showing each species as a different color

2. Create a distance matrix between the 5 canine groups
3. Use principal components analysis to investigate the relationships between the species on the basis of these variables
4. Carry out cluster analysis to study relation between different species.
  - a. Who is Indian Wolf related to?
5. Identify the important factors underlying the Skull measurement
  - a. Is there a relationships between the species with respect to these factors?
6. Carry out a discriminant function analysis to see how well it is possible to separate the groups using the measurements.
7. investigate each canine group separately to see whether logistic regression shows a significant difference between males and females for the measurements. Note that in view of the small sample sizes available for each group, it is unreasonable to expect to fit a logistic function involving all nine variables with good estimates of parameters. Therefore, consideration should be given to fitting functions using only a subset of the variables.
8. Show ROC containing both your discriminant and logistic function for gender classification for the Prehistoric Thai Dog
9. Predict the Gender for the Prehistoric Thai Dog
  - a. Explain the reason for choosing the MVA technique for prediction
  - b. What is the Hit Ratio (Accuracy) of your classification technique?
10. Create a model to predict length of the Mandible length for Prehistoric Thai Dog.
  - a. What is the accuracy of your model

Loading required libraries:

```
library(cluster)
library(data.table)
library(magrittr)
library(stringr)
library(ggplot2)
library(knitr)
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.0 --
```

```
## v tibble  3.0.3    v purrr  0.3.4
## v tidyr   1.1.2    v dplyr  1.0.2
## v readr   1.3.1    v forcats 0.5.0
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::between() masks data.table::between()
## x tidyr::extract() masks magrittr::extract()
## x dplyr::filter() masks stats::filter()
## x dplyr::first() masks data.table::first()
## x dplyr::lag() masks stats::lag()
## x dplyr::last() masks data.table::last()
## x purrr::set_names() masks magrittr::set_names()
## x purrr::transpose() masks data.table::transpose()
```

```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(psych)
```

```
##
```

```
## Attaching package: 'psych'
```

```
## The following objects are masked from 'package:ggplot2':
```

```
##
```

```
## %+%, alpha
```

```
library(FactoMineR)
```

```
library(nFactors)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'nFactors'
```

```
## The following object is masked from 'package:lattice':
```

```
##
```

```
## parallel
```

```
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
```

```
## method from
```

```
## +.gg ggplot2
```

```
library(MASS)
```

```
##  
## Attaching package: 'MASS'  
  
## The following object is masked from 'package:dplyr':  
##  
##     select
```

```
library(gvlma)  
library(leaps)  
library(relaimpo)
```

```
## Loading required package: boot  
  
##  
## Attaching package: 'boot'  
  
## The following object is masked from 'package:lattice':  
##  
##     melanoma  
  
## The following object is masked from 'package:psych':  
##  
##     logit
```

```
## Loading required package: survey
```

```
## Loading required package: grid
```

```
## Loading required package: Matrix
```

```
##  
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':  
##  
##     expand, pack, unpack
```

```
## Loading required package: survival
```

```
##  
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:boot':  
##  
##     aml
```

```

##
## Attaching package: 'survey'

## The following object is masked from 'package:graphics':
##
##     dotchart

## Loading required package: mitools

## This is the global version of package relaimpo.

## If you are a non-US user, a version with the interesting additional metric pmvd is available

## from Ulrike Groempings web site at prof.beuth-hochschule.de/groemping.

library(cowplot)
library(regclass)

## Loading required package: bestglm

## Loading required package: VGAM

## Loading required package: stats4

## Loading required package: splines

##
## Attaching package: 'VGAM'

## The following object is masked from 'package:survey':
##
##     calibrate

## The following objects are masked from 'package:boot':
##
##     logit, simplex

## The following objects are masked from 'package:psych':
##
##     fisherz, logistic, logit

## The following object is masked from 'package:tidyr':
##
##     fill

## Loading required package: rpart

## Loading required package: randomForest

```

```

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:psych':
##
##      outlier

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

## Important regclass change from 1.3:
## All functions that had a . in the name now have an _
## all.correlations -> all_correlations, cor.demo -> cor_demo, etc.

##
## Attaching package: 'regclass'

## The following object is masked from 'package:lattice':
##
##      qq

library(e1071)
library(caret)

##
## Attaching package: 'caret'

## The following object is masked from 'package:VGAM':
##
##      predictors

## The following object is masked from 'package:survival':
##
##      cluster

## The following object is masked from 'package:purrr':
##
##      lift

```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
library(ROCR)
```

```
library(klaR)
```

```
library(readxl)
```

```
library(rmarkdown)
```

```
library(car)
```

```
## Loading required package: carData
```

```
##
```

```
## Attaching package: 'car'
```

```
## The following object is masked from 'package:VGAM':
```

```
##
```

```
##      logit
```

```
## The following object is masked from 'package:boot':
```

```
##
```

```
##      logit
```

```
## The following object is masked from 'package:psych':
```

```
##
```

```
##      logit
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      recode
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      some
```

```
library(SciViews)
```

```
library(psych)
```

```
library(memisc)
```

```
##
```

```
## Attaching package: 'memisc'
```

```
## The following object is masked from 'package:car':
##
##   recode

## The following object is masked from 'package:VGAM':
##
##   Max

## The following object is masked from 'package:Matrix':
##
##   as.array

## The following objects are masked from 'package:dplyr':
##
##   collect, recode, rename, syms

## The following object is masked from 'package:purrr':
##
##   %@%

## The following object is masked from 'package:tibble':
##
##   view

## The following object is masked from 'package:ggplot2':
##
##   syms

## The following objects are masked from 'package:stats':
##
##   contr.sum, contr.treatment, contrasts

## The following object is masked from 'package:base':
##
##   as.array
```

```
#library(FFally)
```

Loading the excel file:

```
df <- read_xlsx("/Users/richardbritto/Desktop/MVA/Final_Data.xlsx")
head(df)
```

```
## # A tibble: 6 x 11
##   CanineGroup    X1    X2    X3    X4    X5    X6    X7    X8    X9 Gender
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
## 1 ModernDog    123  10.1   23   23   19   7.8   32   33   5.6 Male
## 2 ModernDog    137   9.6   19   22   19   7.8   32   40   5.8 Male
## 3 ModernDog    121  10.2   18   21   21   7.9   35   38   6.2 Male
## 4 ModernDog    130  10.7   24   22   20   7.9   32   37   5.9 Male
## 5 ModernDog    149  12     25   25   21   8.4   35   43   6.6 Male
## 6 ModernDog    125   9.5   23   20   20   7.8   33   37   6.3 Male
```



```
attach(df)
```

Exploring the data:

```
dim(df)
```

```
## [1] 77 11
```

```
head(df)
```

```
## # A tibble: 6 x 11
##   CanineGroup    X1    X2    X3    X4    X5    X6    X7    X8    X9 Gender
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
## 1 ModernDog    123  10.1  23   23   19   7.8   32   33   5.6 Male
## 2 ModernDog    137   9.6   19   22   19   7.8   32   40   5.8 Male
## 3 ModernDog    121  10.2   18   21   21   7.9   35   38   6.2 Male
## 4 ModernDog    130  10.7   24   22   20   7.9   32   37   5.9 Male
## 5 ModernDog    149  12     25   25   21   8.4   35   43   6.6 Male
## 6 ModernDog    125   9.5   23   20   20   7.8   33   37   6.3 Male
```

```
tail(df)
```

```
## # A tibble: 6 x 11
##   CanineGroup    X1    X2    X3    X4    X5    X6    X7    X8    X9 Gender
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
## 1 IndianWolves  131  11.8   20   24   23   8.8   38   40   6.5 Female
## 2 IndianWolves  163  10.8   27   24   24   9.2   39   48   7   Female
## 3 IndianWolves  164  10.7   24   23   26   9.5   43   47   7.6 Female
## 4 IndianWolves  141  10.4   20   23   23   8.9   38   43   6   Female
## 5 IndianWolves  148  10.6   26   21   24   8.9   39   40   7   Female
## 6 IndianWolves  158  10.7   25   25   24   9.8   41   45   7.4 Female
```

```
str(df)
```

```
## tibble [77 x 11] (S3: tbl_df/tbl/data.frame)
##  $ CanineGroup: chr [1:77] "ModernDog" "ModernDog" "ModernDog" "ModernDog" ...
##  $ X1         : num [1:77] 123 137 121 130 149 125 126 125 121 122 ...
##  $ X2         : num [1:77] 10.1 9.6 10.2 10.7 12 9.5 9.1 9.7 9.6 8.9 ...
##  $ X3         : num [1:77] 23 19 18 24 25 23 20 19 22 20 ...
##  $ X4         : num [1:77] 23 22 21 22 25 20 22 19 20 20 ...
##  $ X5         : num [1:77] 19 19 21 20 21 20 19 19 18 19 ...
##  $ X6         : num [1:77] 7.8 7.8 7.9 7.9 8.4 7.8 7.5 7.5 7.6 7.6 ...
##  $ X7         : num [1:77] 32 32 35 32 35 33 32 32 31 31 ...
##  $ X8         : num [1:77] 33 40 38 37 43 37 35 37 35 35 ...
##  $ X9         : num [1:77] 5.6 5.8 6.2 5.9 6.6 6.3 5.5 6.2 5.3 5.7 ...
##  $ Gender     : chr [1:77] "Male" "Male" "Male" "Male" ...
```

Cleaning the data:

```
df$CanineGroup <- as.factor(df$CanineGroup)
str(df)
```

```
## tibble [77 x 11] (S3: tbl_df/tbl/data.frame)
## $ CanineGroup: Factor w/ 5 levels "Cuons","GoldenJackal",...: 4 4 4 4 4 4 4 4 4 4 ...
## $ X1         : num [1:77] 123 137 121 130 149 125 126 125 121 122 ...
## $ X2         : num [1:77] 10.1 9.6 10.2 10.7 12 9.5 9.1 9.7 9.6 8.9 ...
## $ X3         : num [1:77] 23 19 18 24 25 23 20 19 22 20 ...
## $ X4         : num [1:77] 23 22 21 22 25 20 22 19 20 20 ...
## $ X5         : num [1:77] 19 19 21 20 21 20 19 19 18 19 ...
## $ X6         : num [1:77] 7.8 7.8 7.9 7.9 8.4 7.8 7.5 7.5 7.6 7.6 ...
## $ X7         : num [1:77] 32 32 35 32 35 33 32 32 31 31 ...
## $ X8         : num [1:77] 33 40 38 37 43 37 35 37 35 35 ...
## $ X9         : num [1:77] 5.6 5.8 6.2 5.9 6.6 6.3 5.5 6.2 5.3 5.7 ...
## $ Gender     : chr [1:77] "Male" "Male" "Male" "Male" ...
```

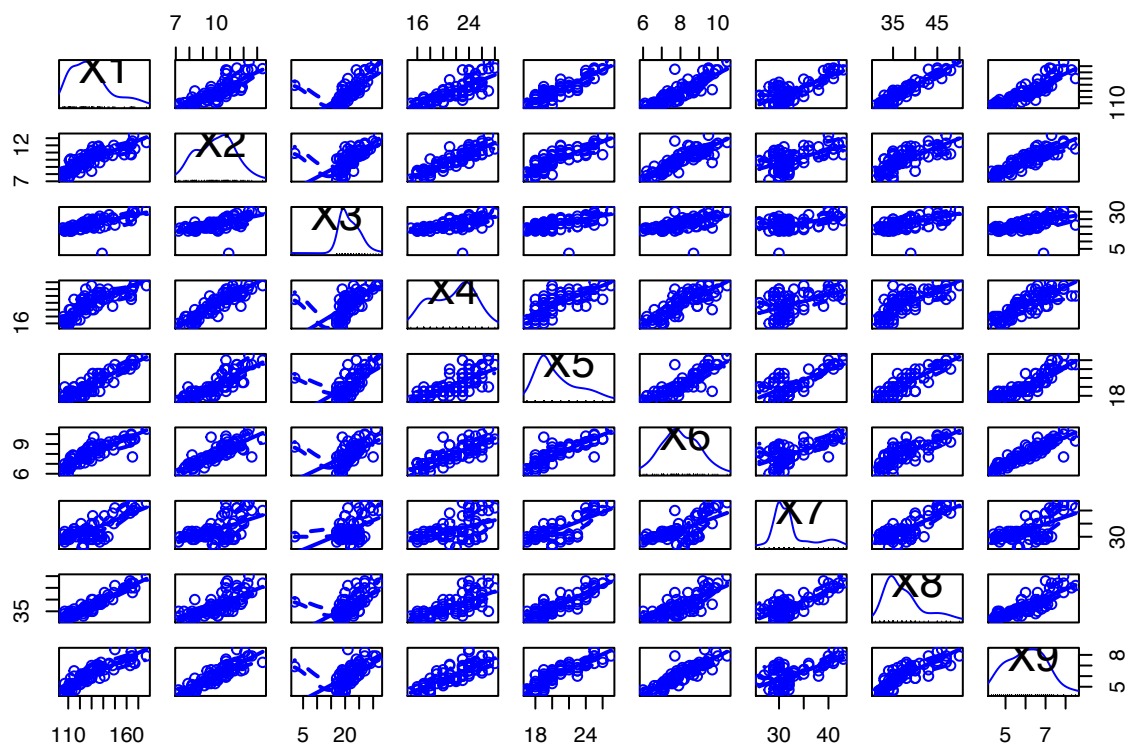
```
df$Gender <- as.factor(df$Gender)
str(df)
```

```
## tibble [77 x 11] (S3: tbl_df/tbl/data.frame)
## $ CanineGroup: Factor w/ 5 levels "Cuons","GoldenJackal",...: 4 4 4 4 4 4 4 4 4 4 ...
## $ X1         : num [1:77] 123 137 121 130 149 125 126 125 121 122 ...
## $ X2         : num [1:77] 10.1 9.6 10.2 10.7 12 9.5 9.1 9.7 9.6 8.9 ...
## $ X3         : num [1:77] 23 19 18 24 25 23 20 19 22 20 ...
## $ X4         : num [1:77] 23 22 21 22 25 20 22 19 20 20 ...
## $ X5         : num [1:77] 19 19 21 20 21 20 19 19 18 19 ...
## $ X6         : num [1:77] 7.8 7.8 7.9 7.9 8.4 7.8 7.5 7.5 7.6 7.6 ...
## $ X7         : num [1:77] 32 32 35 32 35 33 32 32 31 31 ...
## $ X8         : num [1:77] 33 40 38 37 43 37 35 37 35 35 ...
## $ X9         : num [1:77] 5.6 5.8 6.2 5.9 6.6 6.3 5.5 6.2 5.3 5.7 ...
## $ Gender     : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 1 1 ...
```

1. Using suitable graphical method, compare the distribution of the nine variables for the prehistoric and modern Thai dog.

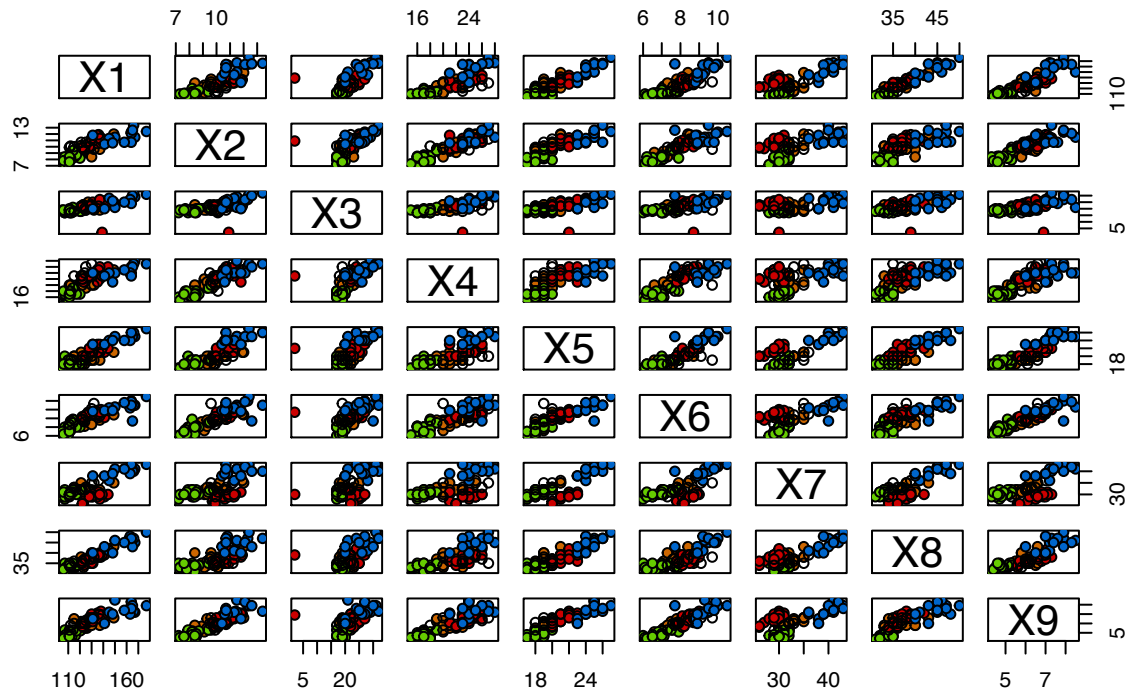
a. Create a Draftsman plot for the 9 variables showing each species as a different color

```
scatterplotMatrix(df[, 2:10])
```




```
pairs(df[2:10],
      main = "Skull Measurements for all Canine species",
      pch = 21,
      bg = c('#CC0000', '#66CC00', '#0066CC', '#CC6600')
      [unclass(df$CanineGroup)])
```

## Skull Measurements for all Canine species



2. Create a distance matrix between the 5 canine groups:

```
mat.std = scale(df[c(-1, -11)])
dist.df.scaled <- dist(mat.std, method = "euclidean")
plot(dist.df.scaled)
```



3. Use principal components analysis to investigate the relationships between the species on the basis of these variables

```
cor(df[c(-1, -11)])
```

```
##           X1           X2           X3           X4           X5           X6           X7
## X1 1.0000000 0.8259623 0.6841756 0.7976348 0.9066471 0.8515578 0.7589012
## X2 0.8259623 1.0000000 0.6184360 0.8897336 0.8213389 0.8457847 0.5597767
## X3 0.6841756 0.6184360 1.0000000 0.6200059 0.6166557 0.5608910 0.4516023
## X4 0.7976348 0.8897336 0.6200059 1.0000000 0.7402734 0.8085781 0.4707245
## X5 0.9066471 0.8213389 0.6166557 0.7402734 1.0000000 0.8537794 0.7424201
## X6 0.8515578 0.8457847 0.5608910 0.8085781 0.8537794 1.0000000 0.6456683
## X7 0.7589012 0.5597767 0.4516023 0.4707245 0.7424201 0.6456683 1.0000000
## X8 0.9494620 0.7460676 0.5906419 0.7151408 0.8777774 0.7984086 0.7867110
## X9 0.8833714 0.8874866 0.5750451 0.8229495 0.8826925 0.8942284 0.6478342
##           X8           X9
## X1 0.9494620 0.8833714
## X2 0.7460676 0.8874866
## X3 0.5906419 0.5750451
## X4 0.7151408 0.8229495
## X5 0.8777774 0.8826925
## X6 0.7984086 0.8942284
## X7 0.7867110 0.6478342
## X8 1.0000000 0.8380353
## X9 0.8380353 1.0000000
```

```
data_pca <- prcomp(df[, -c(1, 11)], scale = TRUE)
```

```
temp = as.matrix(df[, 1])  
class(temp)
```

```
## [1] "matrix"
```

```
temp
```

```
##      CanineGroup  
## [1,] "ModernDog"  
## [2,] "ModernDog"  
## [3,] "ModernDog"  
## [4,] "ModernDog"  
## [5,] "ModernDog"  
## [6,] "ModernDog"  
## [7,] "ModernDog"  
## [8,] "ModernDog"  
## [9,] "ModernDog"  
## [10,] "ModernDog"  
## [11,] "ModernDog"  
## [12,] "ModernDog"  
## [13,] "ModernDog"  
## [14,] "ModernDog"  
## [15,] "ModernDog"  
## [16,] "ModernDog"  
## [17,] "GoldenJackal"  
## [18,] "GoldenJackal"  
## [19,] "GoldenJackal"  
## [20,] "GoldenJackal"  
## [21,] "GoldenJackal"  
## [22,] "GoldenJackal"  
## [23,] "GoldenJackal"  
## [24,] "GoldenJackal"  
## [25,] "GoldenJackal"  
## [26,] "GoldenJackal"  
## [27,] "GoldenJackal"  
## [28,] "GoldenJackal"  
## [29,] "GoldenJackal"  
## [30,] "GoldenJackal"  
## [31,] "GoldenJackal"  
## [32,] "GoldenJackal"  
## [33,] "GoldenJackal"  
## [34,] "GoldenJackal"  
## [35,] "GoldenJackal"  
## [36,] "GoldenJackal"  
## [37,] "Cuons"  
## [38,] "Cuons"  
## [39,] "Cuons"  
## [40,] "Cuons"  
## [41,] "Cuons"  
## [42,] "Cuons"
```

```

## [43,] "Cuons"
## [44,] "Cuons"
## [45,] "Cuons"
## [46,] "Cuons"
## [47,] "Cuons"
## [48,] "Cuons"
## [49,] "Cuons"
## [50,] "Cuons"
## [51,] "Cuons"
## [52,] "Cuons"
## [53,] "Cuons"
## [54,] "ThaiDogs"
## [55,] "ThaiDogs"
## [56,] "ThaiDogs"
## [57,] "ThaiDogs"
## [58,] "ThaiDogs"
## [59,] "ThaiDogs"
## [60,] "ThaiDogs"
## [61,] "ThaiDogs"
## [62,] "ThaiDogs"
## [63,] "ThaiDogs"
## [64,] "IndianWolves"
## [65,] "IndianWolves"
## [66,] "IndianWolves"
## [67,] "IndianWolves"
## [68,] "IndianWolves"
## [69,] "IndianWolves"
## [70,] "IndianWolves"
## [71,] "IndianWolves"
## [72,] "IndianWolves"
## [73,] "IndianWolves"
## [74,] "IndianWolves"
## [75,] "IndianWolves"
## [76,] "IndianWolves"
## [77,] "IndianWolves"

```

data\_pca

```

## Standard deviations (1, ..., p=9):
## [1] 2.6555963 0.8391652 0.7365758 0.4390554 0.4241988 0.3627806 0.3031519
## [8] 0.2652189 0.1857418
##
## Rotation (n x k) = (9 x 9):
##          PC1          PC2          PC3          PC4          PC5          PC6
## X1 0.3636408 -0.11451510 0.08210471 -0.30326354 0.24950692 -0.07899550
## X2 0.3424554 0.31490128 -0.19979188 0.33605928 0.01517931 0.49451257
## X3 0.2665621 0.32018675 0.87894338 0.04161625 -0.18169514 -0.04568559
## X4 0.3265349 0.44638084 -0.16540131 0.26534253 0.54545187 -0.21526217
## X5 0.3539586 -0.14160855 -0.03861441 -0.26352534 -0.33012092 0.43239890
## X6 0.3459444 0.06792334 -0.26250857 0.05378069 -0.51974026 -0.68294862
## X7 0.2859405 -0.68736531 0.13651981 0.64014932 0.05443187 -0.01170970
## X8 0.3470802 -0.28877388 0.03666665 -0.47256682 0.40753260 -0.10978603
## X9 0.3544268 0.07362113 -0.25111557 -0.13231892 -0.24254817 0.18765482
##          PC7          PC8          PC9

```

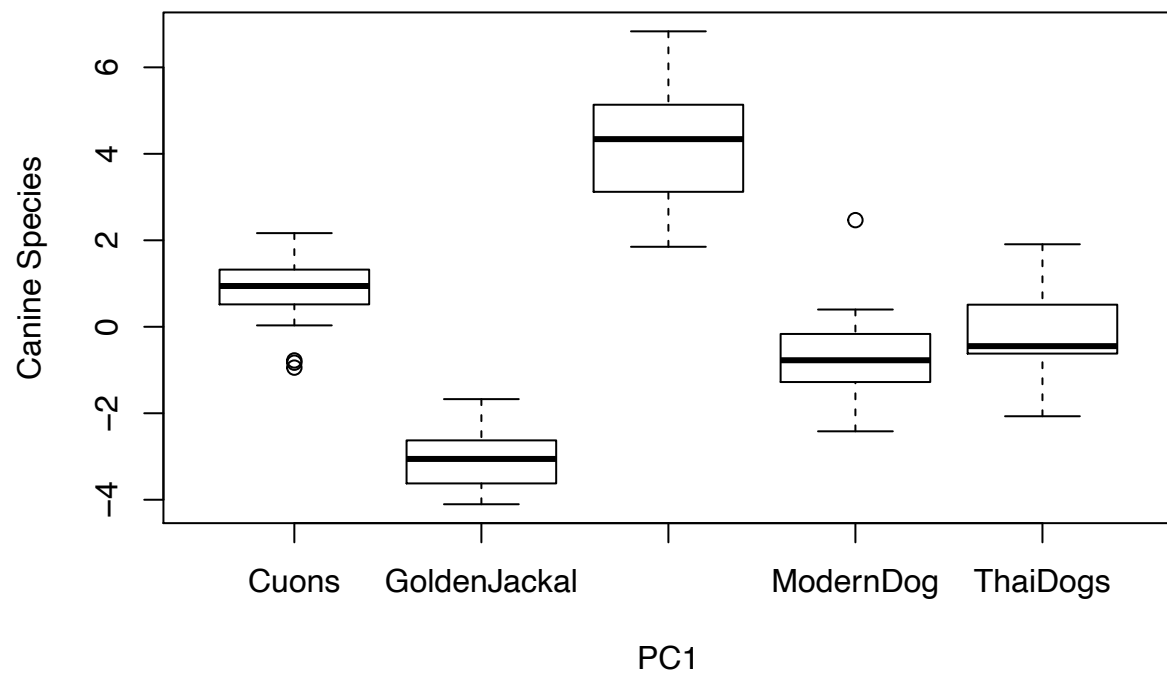
```
## X1  0.05543869  0.16005914  0.811637429
## X2  0.13657790  0.60411640 -0.048224206
## X3  0.08257828 -0.03476461 -0.094992855
## X4 -0.30849700 -0.39244126 -0.057608736
## X5 -0.67024916 -0.19081879 -0.046144083
## X6 -0.08734948  0.23986950 -0.046794522
## X7  0.03463451 -0.11415807  0.002559605
## X8  0.12889717  0.23397418 -0.567438948
## X9  0.63372357 -0.54081471 -0.016253801
```


```
predict(data_pca)[, 1]
```


```
## [1] -0.68594144 -0.24567287 -0.08646546  0.16794155  2.46605602 -0.31165997
## [7] -1.13265883 -0.95932707 -1.42772439 -1.42312985 -1.08479346 -2.41795380
## [13] -1.49439221 -0.59855768 -0.86184015  0.39898643 -2.50221460 -2.66854242
## [19] -3.13382930 -2.58652273 -2.21301819 -2.97976934 -2.78947778 -2.53185477
## [25] -1.67354577 -2.78123568 -2.82501243 -3.30622428 -4.10431266 -3.66153536
## [31] -4.04399515 -3.72261526 -3.82745453 -3.58415662 -3.18995154 -3.28698703
## [37] -0.83046450  1.69865955  1.94265070  1.24244782  1.32210362  1.18642807
## [43]  0.94267205  0.83611277  1.18696596  0.58012526  0.87381349  2.16615489
## [49]  0.03166471 -0.94467005  1.96212575  0.51806252 -0.77814254 -1.81012401
## [55] -0.60672494  1.90795558 -2.06926137  1.40688395  0.51098957  0.42526384
## [61] -0.37712553 -0.62016146 -0.51807155  4.97863584  5.43014384  4.19582141
## [67]  3.12145386  6.83250997  5.78067966  5.13564203  4.48437886  1.98423294
## [73]  3.99779940  4.49728950  1.85001891  2.55231795  4.08213029
```


```
out <- sapply(1:5, function(i){plot(df$CanineGroup,
                                   data_pca$x[, i],
                                   xlab=paste("PC", i, sep = ""),
                                   ylab="Canine Species")})
```










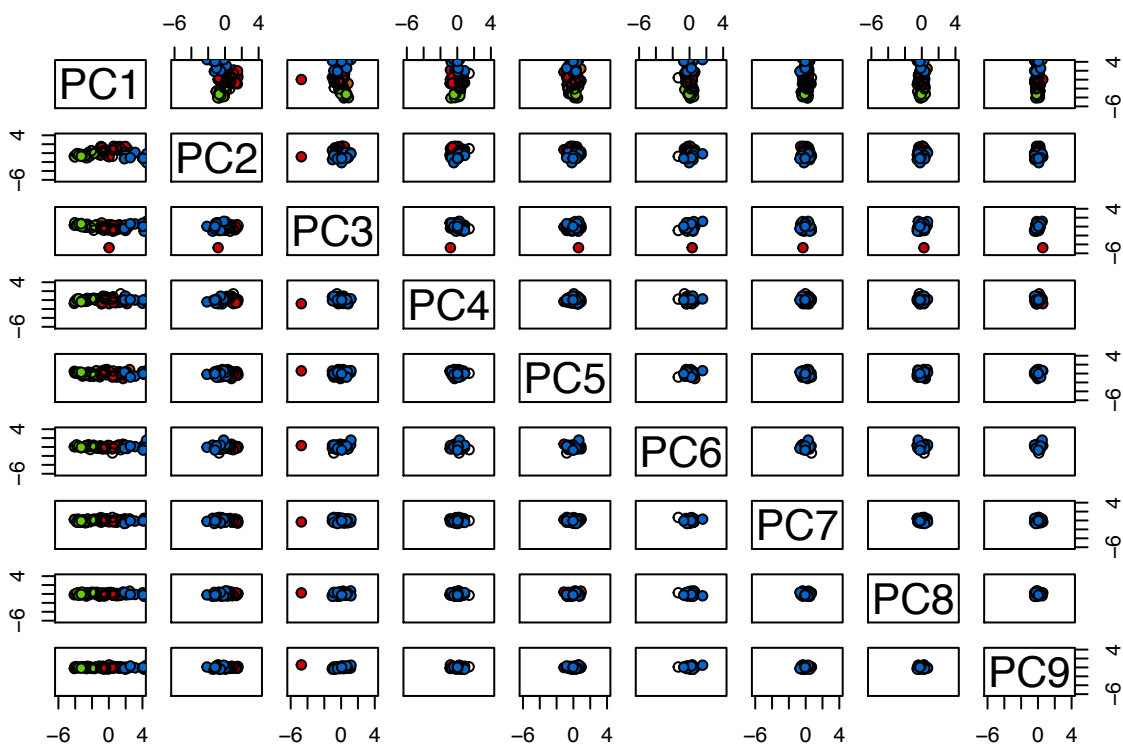




Below we can see the relation between each variables for each species:

```
clr <- character(nrow(df))
clr[] <- "black"
clr[df$CanineGroup == "ModernDog"] <- "#CC0000"
clr[df$CanineGroup == "GoldenJackal"] <- "#66CC00"
clr[df$CanineGroup == "Cuons"] <- "#0066CC"
clr[df$CanineGroup == "ThaiDogs"] <- "#CC6600"
clr[df$CanineGroup == "IndianWolves"] <- "#6600CC"

pairs(data_pca$x[, 1:9],
      ylim = c(-6, 4),
      xlim = c(-6, 4),
      pch = 21,
      bg = c('#CC0000', '#66CC00', '#0066CC', '#CC6600')
      [unclass(df$CanineGroup)])
```



```
summary(data_pca)
```

```
## Importance of components:
##               PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  2.6556 0.83917 0.73658 0.43906 0.42420 0.36278 0.30315
## Proportion of Variance 0.7836 0.07824 0.06028 0.02142 0.01999 0.01462 0.01021
## Cumulative Proportion 0.7836 0.86182 0.92210 0.94352 0.96352 0.97814 0.98835
##               PC8      PC9
## Standard deviation  0.26522 0.18574
## Proportion of Variance 0.00782 0.00383
## Cumulative Proportion 0.99617 1.00000
```

loadings (eigenvectors) are stored in `data_pca$rotation`:

```
data_pca$rotation
```

```
##           PC1           PC2           PC3           PC4           PC5           PC6
## X1  0.3636408 -0.11451510  0.08210471 -0.30326354  0.24950692 -0.07899550
## X2  0.3424554  0.31490128 -0.19979188  0.33605928  0.01517931  0.49451257
## X3  0.2665621  0.32018675  0.87894338  0.04161625 -0.18169514 -0.04568559
## X4  0.3265349  0.44638084 -0.16540131  0.26534253  0.54545187 -0.21526217
## X5  0.3539586 -0.14160855 -0.03861441 -0.26352534 -0.33012092  0.43239890
## X6  0.3459444  0.06792334 -0.26250857  0.05378069 -0.51974026 -0.68294862
## X7  0.2859405 -0.68736531  0.13651981  0.64014932  0.05443187 -0.01170970
```

```
## X8 0.3470802 -0.28877388 0.03666665 -0.47256682 0.40753260 -0.10978603
## X9 0.3544268 0.07362113 -0.25111557 -0.13231892 -0.24254817 0.18765482
##          PC7          PC8          PC9
## X1 0.05543869 0.16005914 0.811637429
## X2 0.13657790 0.60411640 -0.048224206
## X3 0.08257828 -0.03476461 -0.094992855
## X4 -0.30849700 -0.39244126 -0.057608736
## X5 -0.67024916 -0.19081879 -0.046144083
## X6 -0.08734948 0.23986950 -0.046794522
## X7 0.03463451 -0.11415807 0.002559605
## X8 0.12889717 0.23397418 -0.567438948
## X9 0.63372357 -0.54081471 -0.016253801
```

Eigenvalues are  $sdev^2$ :

```
eigen_data<-data_pca$sdev^2
```


Sample scores stored in `data_pca$x`:

```
head(data_pca$x)
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6
## [1,] -0.68594144 0.7845059 0.3006808 0.87047017 0.099293422 -0.13703394
## [2,] -0.24567287 -0.2998432 -0.3386984 -0.38688425 0.904531523 -0.40703561
## [3,] -0.08646546 -0.7099032 -0.7324535 0.36215960 0.007597324 0.34684409
## [4,] 0.16794155 0.5263915 0.4238685 0.25556873 0.116376021 0.15817439
## [5,] 2.46605602 0.3055451 0.2215356 0.12938866 0.883927889 0.13952874
## [6,] -0.31165997 -0.1929752 0.4205179 -0.01594082 -0.279041559 0.03359648
##          PC7          PC8          PC9
## [1,] -0.12574305 -0.08103187 0.27299561
## [2,] 0.21201705 0.24663975 0.14167354
## [3,] -0.04042272 -0.04720248 -0.36986352
## [4,] 0.09155686 0.34942670 0.02832178
## [5,] 0.34760207 0.63208273 -0.03260669
## [6,] 0.38738491 -0.23534059 -0.10684702
```

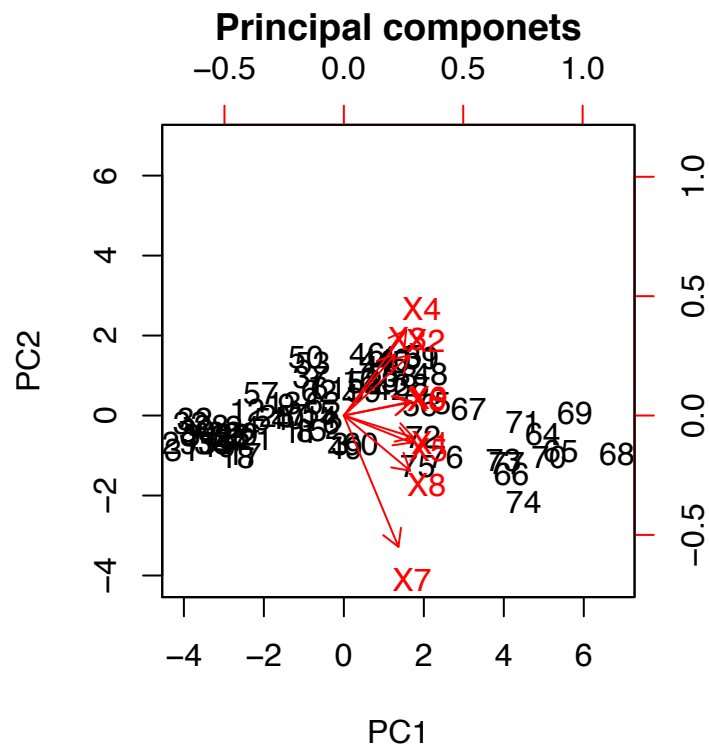
Viewing the scree plot to see how many PCs are optimal for our data and we can conclude that maximum variance is covered by the top 3 PCs i.e the PC1, PC2, PC3:

```
fviz_eig(data_pca)
```




```
biplot(data_pca, scale = 0, main = 'Principal componets')
```





We can see the contributions of each variables to the top 2 dimensions and this gives us some good insights on how do the variables contributre to all the data points in our data:

```
fviz_pca_var(data_pca,
              col.var = "contrib",
              gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
              repel = TRUE)
```



4. Carry out cluster analysis to study relation between different species.

a. Who is Indian Wolf related to?

Hierarchical cluster analysis, Nearest-neighbor Standardizing the data with scale()

```
temp_values = c('X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X9')
df_species = df[, -11]
df_species
```

```
## # A tibble: 77 x 10
##   CanineGroup   X1    X2    X3    X4    X5    X6    X7    X8    X9
##   <fct>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 ModernDog   123  10.1  23   23   19   7.8   32   33   5.6
## 2 ModernDog   137   9.6  19   22   19   7.8   32   40   5.8
## 3 ModernDog   121  10.2  18   21   21   7.9   35   38   6.2
## 4 ModernDog   130  10.7  24   22   20   7.9   32   37   5.9
## 5 ModernDog   149  12    25   25   21   8.4   35   43   6.6
## 6 ModernDog   125   9.5  23   20   20   7.8   33   37   6.3
## 7 ModernDog   126   9.1  20   22   19   7.5   32   35   5.5
## 8 ModernDog   125   9.7  19   19   19   7.5   32   37   6.2
## 9 ModernDog   121   9.6  22   20   18   7.6   31   35   5.3
## 10 ModernDog  122   8.9  20   20   19   7.6   31   35   5.7
## # ... with 67 more rows
```

```
df_species = df_species %>%
  group_by(CanineGroup) %>%
  summarise_at(temp_values, mean, na.rm = TRUE)
```

```
setDT(df_species)
```

```
df_species = df_species %>%
  remove_rownames %>%
  column_to_rownames(var = "CanineGroup")
```

```
mat.std_species = scale(df_species)
```

Creating a (Euclidean) distance matrix of the standardized data :

```
dist_species <- dist(mat.std_species, method = "euclidean")
```


```
clusspecies_nn <- hclust(dist_species, method = "single")
```

```
dist_species
```

```
##              Cuons GoldenJackal IndianWolves ModernDog
## GoldenJackal 4.680179
## IndianWolves 4.059156      7.806678
## ModernDog    2.062237      2.831293      5.171149
## ThaiDogs     1.818116      3.524634      4.964169      1.137355
```

As we see above that Indian Wolf is closed to Cuans with a distance of 4.059156. Futher, we can visualize this using dendrogram or a kmeans plot.


```
# For Nearest Neighbor - Single Linkage:
par(mar = c(8, 4, 2, 1) + 0.1)
plot(as.dendrogram(clusspecies_nn),
     ylab = "Distance between Canine Species",
     ylim = c(0, 5.5),
     main = "Dendrogram - Nearest Neighbor")
```



IndianWolves is closest to Cuons

#### Kmeans

```
kmeans_3_species <- kmeans(mat.std_species, 3, nstart = 10)
fviz_cluster(kmeans_3_species, data = df_species)
```



## 5. Identify the important factors underlying the Skull measurement

### a. Is there a relationships between the species with respect to these factors?


Calculating the correlation matrix for all the numeric data in our dataset

```
corr.mat = cor(df[, 2:10])
corr.mat
```

```
##          X1          X2          X3          X4          X5          X6          X7
## X1 1.0000000 0.8259623 0.6841756 0.7976348 0.9066471 0.8515578 0.7589012
## X2 0.8259623 1.0000000 0.6184360 0.8897336 0.8213389 0.8457847 0.5597767
## X3 0.6841756 0.6184360 1.0000000 0.6200059 0.6166557 0.5608910 0.4516023
## X4 0.7976348 0.8897336 0.6200059 1.0000000 0.7402734 0.8085781 0.4707245
## X5 0.9066471 0.8213389 0.6166557 0.7402734 1.0000000 0.8537794 0.7424201
## X6 0.8515578 0.8457847 0.5608910 0.8085781 0.8537794 1.0000000 0.6456683
## X7 0.7589012 0.5597767 0.4516023 0.4707245 0.7424201 0.6456683 1.0000000
## X8 0.9494620 0.7460676 0.5906419 0.7151408 0.8777774 0.7984086 0.7867110
## X9 0.8833714 0.8874866 0.5750451 0.8229495 0.8826925 0.8942284 0.6478342
##          X8          X9
## X1 0.9494620 0.8833714
## X2 0.7460676 0.8874866
## X3 0.5906419 0.5750451
## X4 0.7151408 0.8229495
## X5 0.8777774 0.8826925
```

```
## X6 0.7984086 0.8942284
## X7 0.7867110 0.6478342
## X8 1.0000000 0.8380353
## X9 0.8380353 1.0000000
```

```
plot(corr.mat)
```



This is the correlation plot


Calculating the PCA and plotting these variances:

```
df_pca <- prcomp(df[2:10], scale = TRUE)
summary(df_pca)
```

```
## Importance of components:
```

```
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  2.6556 0.83917 0.73658 0.43906 0.42420 0.36278 0.30315
## Proportion of Variance 0.7836 0.07824 0.06028 0.02142 0.01999 0.01462 0.01021
## Cumulative Proportion 0.7836 0.86182 0.92210 0.94352 0.96352 0.97814 0.98835
##          PC8      PC9
## Standard deviation  0.26522 0.18574
## Proportion of Variance 0.00782 0.00383
## Cumulative Proportion 0.99617 1.00000
```

```
plot(df_pca)
```



As from Q3 we can see that the top 3 PCs restore maximum variance

A table containing eigenvalues and %'s accounted, follows. Eigenvalues are the  $sdev^2$ :

```
eigen_values <- round(df_pca$sdev^2, 2)
head(eigen_values)
```

```
## [1] 7.05 0.70 0.54 0.19 0.18 0.13
```

Eigen vectors as follows:

```
eigen_vectors <- df_pca$rotation
```

Taking the 3 PCs to generate linear combinations for all the variables:

```
pcafactors <- eigen_vectors[, 1:3]
pcafactors
```

```
##          PC1          PC2          PC3
## X1 0.3636408 -0.11451510 0.08210471
## X2 0.3424554 0.31490128 -0.19979188
## X3 0.2665621 0.32018675 0.87894338
```

```
## X4 0.3265349  0.44638084 -0.16540131
## X5 0.3539586 -0.14160855 -0.03861441
## X6 0.3459444  0.06792334 -0.26250857
## X7 0.2859405 -0.68736531  0.13651981
## X8 0.3470802 -0.28877388  0.03666665
## X9 0.3544268  0.07362113 -0.25111557
```

Multiplying each column of the eigenvector's matrix by the square-root of the corresponding eigenvalue in order to get the factor loadings:

```
unrot.fact <- sweep(pcafactors, MARGIN = 2, df_pca$sdev[1:3], '*')
unrot.fact
```

```
##          PC1          PC2          PC3
## X1 0.9656832 -0.09609709  0.06047634
## X2 0.9094234  0.26425421 -0.14716186
## X3 0.7078814  0.26868959  0.64740843
## X4 0.8671448  0.37458728 -0.12183061
## X5 0.9399711 -0.11883297 -0.02844244
## X6 0.9186886  0.05699890 -0.19335746
## X7 0.7593427 -0.57681308  0.10055719
## X8 0.9217049 -0.24232900  0.02700777
## X9 0.9412145  0.06178029 -0.18496566
```

Computing communalities:

```
communalities <- rowSums(unrot.fact^2)
communalities
```

```
##          X1          X2          X3          X4          X5          X6          X7          X8
## 0.9454361 0.9185378 0.9924279 0.9070985 0.8984760 0.8846248 0.9194263 0.9089926
##          X9
## 0.9239138
```

Performing the varimax rotation. The default in the varimax function is norm=TRUE thus, Kaiser normalization is carried out:

```
rot.fact <- varimax(unrot.fact)
rot.fact
```

```
## $loadings
##
## Loadings:
##    PC1    PC2    PC3
## X1 0.622 -0.652  0.365
## X2 0.858 -0.314  0.289
## X3 0.324 -0.252  0.908
## X4 0.869 -0.202  0.334
## X5 0.636 -0.649  0.270
## X6 0.787 -0.482  0.181
## X7 0.220 -0.920  0.159
## X8 0.536 -0.741  0.269
```



```
## X9  0.802 -0.492  0.198
##
##          PC1    PC2    PC3
## SS loadings  3.985 2.919 1.395
## Proportion Var 0.443 0.324 0.155
## Cumulative Var 0.443 0.767 0.922
##
## $rotmat
##          [,1]      [,2]      [,3]
## [1,]  0.7231780 -0.59076663 0.3577826
## [2,]  0.4847295  0.80315318 0.3463846
## [3,] -0.4919867 -0.07706991 0.8671847
```

The print method of varimax omits loadings less than `abs(0.1)`.

```
fact.load <- rot.fact$loadings[1:3, 1:3]
fact.load
```

```
##          PC1          PC2          PC3
## X1 0.6220262 -0.6523350 0.3646623
## X2 0.8581685 -0.3136786 0.2892930
## X3 0.3236497 -0.2522895 0.9077603
```

Computing the rotated factor scores for the 25 European Countries.

```
scale <- scale(df[, 2:10])
head(scale)
```

```
##          X1          X2          X3          X4          X5          X6
## [1,] -0.34133663  0.09904431  0.3235696  0.4459670 -0.5997769 -0.19537598
## [2,]  0.45857835 -0.25733007 -0.6255679  0.1499372 -0.5997769 -0.19537598
## [3,] -0.45561020  0.17031919 -0.8628523 -0.1460926  0.2034026 -0.09768799
## [4,]  0.05862086  0.52669357  0.5608540  0.1499372 -0.1981871 -0.09768799
## [5,]  1.14421975  1.45326697  0.7981384  1.0380266  0.2034026  0.39075196
## [6,] -0.22706306 -0.32860495  0.3235696 -0.4421225 -0.1981871 -0.19537598
##          X7          X8          X9
## [1,] -0.1244973 -0.99951775 -0.4661440
## [2,] -0.1244973  0.58968599 -0.2700069
## [3,]  0.5944746  0.13562778  0.1222673
## [4,] -0.1244973 -0.09140133 -0.1719384
## [5,]  0.5944746  1.27077331  0.5145414
## [6,]  0.1151600 -0.09140133  0.2203358
```

```
fit.pc <- principal(df[, 2:10], nfactors = 3, rotate = "varimax")
fit.pc
```

```
## Principal Components Analysis
## Call: principal(r = df[, 2:10], nfactors = 3, rotate = "varimax")
## Standardized loadings (pattern matrix) based upon correlation matrix
##      RC1 RC2 RC3 h2      u2 com
## X1 0.62 0.65 0.36 0.95 0.0546 2.6
## X2 0.86 0.31 0.29 0.92 0.0815 1.5
```

```
## X3 0.32 0.25 0.91 0.99 0.0076 1.4
## X4 0.87 0.20 0.33 0.91 0.0929 1.4
## X5 0.64 0.65 0.27 0.90 0.1015 2.3
## X6 0.79 0.48 0.18 0.88 0.1154 1.8
## X7 0.22 0.92 0.16 0.92 0.0806 1.2
## X8 0.54 0.74 0.27 0.91 0.0910 2.1
## X9 0.80 0.49 0.20 0.92 0.0761 1.8
##
##
##          RC1  RC2  RC3
## SS loadings      3.98 2.92 1.40
## Proportion Var    0.44 0.32 0.16
## Cumulative Var    0.44 0.77 0.92
## Proportion Explained 0.48 0.35 0.17
## Cumulative Proportion 0.48 0.83 1.00
##
## Mean item complexity = 1.8
## Test of the hypothesis that 3 components are sufficient.
##
## The root mean square of the residuals (RMSR) is 0.02
## with the empirical chi square 3 with prob < 1
##
## Fit based upon off diagonal values = 1
```

```
round(fit.pc$values, 3)
```

```
## [1] 7.052 0.704 0.543 0.193 0.180 0.132 0.092 0.070 0.035
```

```
fit.pc$loadings
```

```
##
## Loadings:
##      RC1  RC2  RC3
## X1 0.622 0.652 0.365
## X2 0.858 0.314 0.289
## X3 0.324 0.252 0.908
## X4 0.869 0.202 0.334
## X5 0.636 0.649 0.270
## X6 0.787 0.482 0.181
## X7 0.220 0.920 0.159
## X8 0.536 0.741 0.269
## X9 0.802 0.492 0.198
##
##          RC1  RC2  RC3
## SS loadings 3.985 2.919 1.395
## Proportion Var 0.443 0.324 0.155
## Cumulative Var 0.443 0.767 0.922
```

Loadings with more digits

```
for (i in c(3, 1)) { print(fit.pc$loadings[[1, i]])}
```

```
## [1] 0.3646623
## [1] 0.6220262
```

Communalities

```
fit.pc$communality
```

```
##          X1          X2          X3          X4          X5          X6          X7          X8
## 0.9454361 0.9185378 0.9924279 0.9070985 0.8984760 0.8846248 0.9194263 0.9089926
##          X9
## 0.9239138
```


Mandible is able to restore 95% of the total variance Rotated factor scores

```
head(fit.pc$scores)
```

```
##          RC1          RC2          RC3
## [1,] 0.06552327 -0.8719736 0.5854046
## [2,] -0.01387213 0.1968842 -0.5556222
## [3,] 0.05562332 0.5835645 -1.1670095
## [4,] 0.06667753 -0.4220909 0.7389350
## [5,] 0.70008295 0.2793482 0.7191850
## [6,] -0.47722049 0.1593615 0.3734396
```

FA utilities

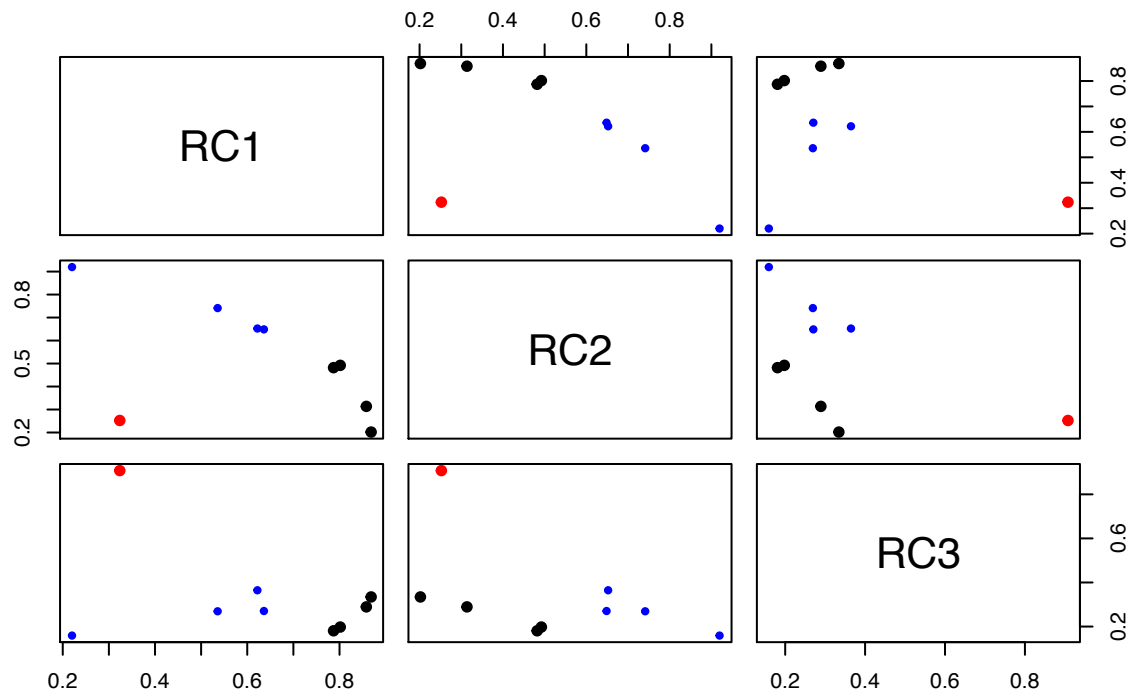
```
fa.parallel(df[, 2:10]) # See factor recommendation
```



```
## Parallel analysis suggests that the number of factors = 1 and the number of components = 1
```

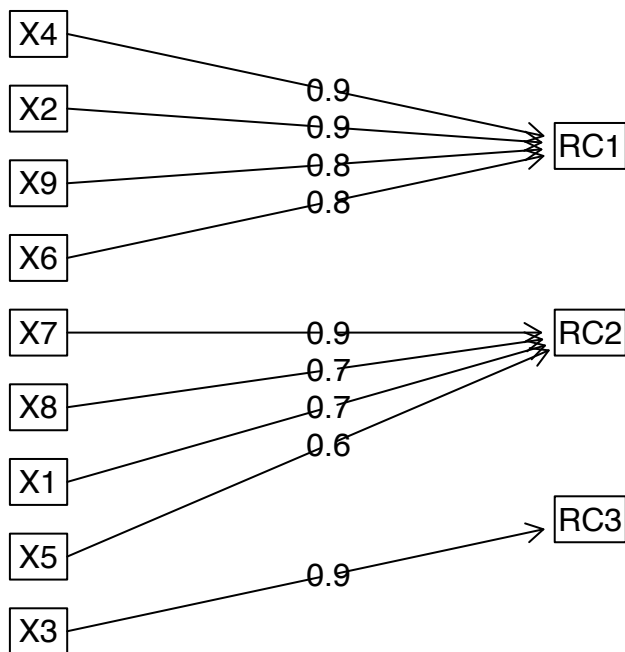
```
fa.plot(fit.pc) # See Correlations within Factors
```

## Principal Component Analysis




```
fa.diagram(fit.pc) # Visualize the relationship
```

## Components Analysis



```
vss(df[, 2:10]) # See Factor recommendations for a simple structure
```

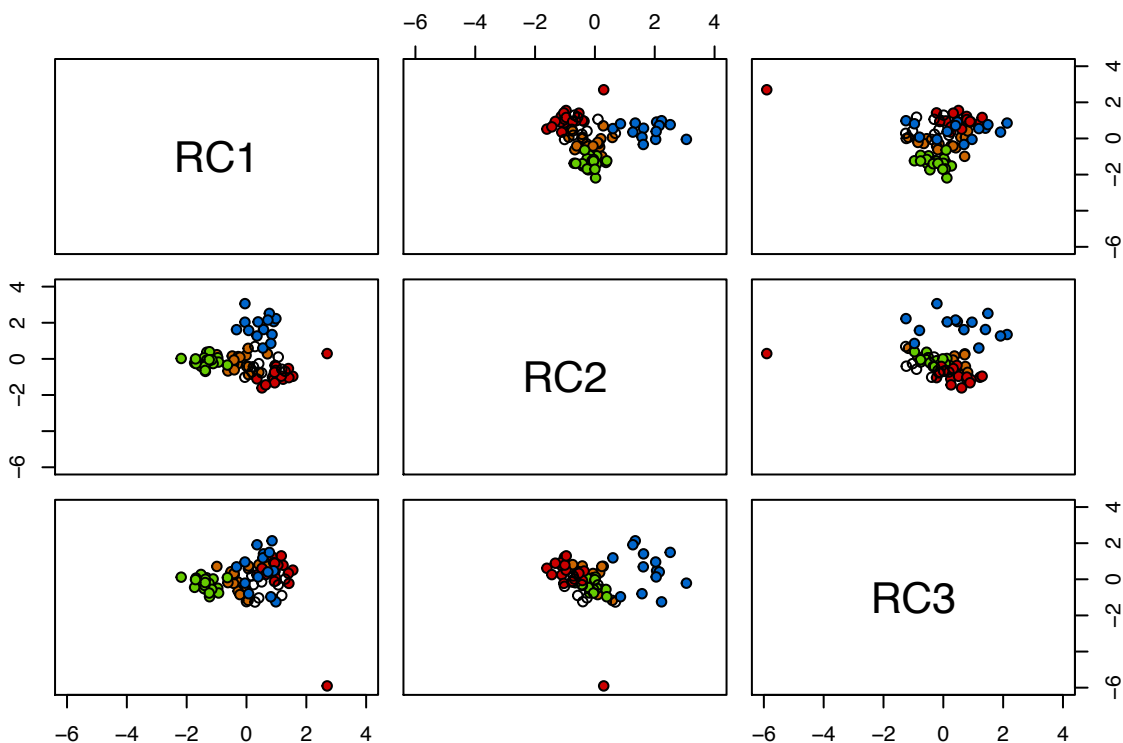


```
##
## Very Simple Structure
## Call: vss(x = df[, 2:10])
## VSS complexity 1 achieves a maximum of 0.98 with 1 factors
## VSS complexity 2 achieves a maximum of 0.99 with 2 factors
##
## The Velicer MAP achieves a minimum of 0.08 with 2 factors
## BIC achieves a minimum of NA with 2 factors
## Sample Size adjusted BIC achieves a minimum of NA with 4 factors
##
## Statistics by number of factors
##   vss1 vss2  map dof   chisq   prob sqresid  fit RMSEA BIC  SABIC complex
## 1 0.98 0.00 0.094 27 1.4e+02 2.6e-17 0.98 0.98 0.233 23 108.37 1.0
## 2 0.51 0.99 0.080 19 4.4e+01 8.6e-04 0.51 0.99 0.131 -38 21.67 1.7
## 3 0.44 0.85 0.121 12 1.5e+01 2.2e-01 0.38 0.99 0.059 -37 1.12 2.1
## 4 0.45 0.92 0.165 6 4.1e+00 6.6e-01 0.23 1.00 0.000 -22 -3.05 2.0
## 5 0.45 0.91 0.174 1 3.0e-01 5.8e-01 0.16 1.00 0.000 -4 -0.89 2.0
## 6 0.44 0.84 0.288 -3 4.7e-08 NA 0.13 1.00 NA NA NA 2.2
## 7 0.44 0.89 0.550 -6 1.4e-06 NA 0.28 0.99 NA NA NA 2.2
## 8 0.44 0.89 1.000 -8 0.0e+00 NA 0.28 0.99 NA NA NA 2.2
##   eChisq SRMR eCRMS eBIC
## 1 1.7e+01 5.5e-02 0.0635 -100.5
## 2 2.1e+00 2.0e-02 0.0269 -80.4
## 3 4.1e-01 8.6e-03 0.0150 -51.7
## 4 9.5e-02 4.1e-03 0.0101 -26.0
## 5 5.3e-03 9.8e-04 0.0059 -4.3
```

```
## 6 8.6e-10 3.9e-07 NA NA
## 7 1.4e-08 1.6e-06 NA NA
## 8 9.9e-18 4.2e-11 NA NA
```

```
clr <- character(nrow(df))
clr[] <- "black"
clr[df$CanineGroup == "ModernDog"] <- "#CC0000"
clr[df$CanineGroup == "GoldenJackal"] <- "#66CC00"
clr[df$CanineGroup == "Cuons"] <- "#0066CC"
clr[df$CanineGroup == "ThaiDogs"] <- "#CC6600"
clr[df$CanineGroup == "IndianWolves"] <- "#6600CC"

pairs(fit.pc$scores[, 1:3],
      ylim = c(-6, 4),
      xlim = c(-6, 4),
      pch = 21,
      bg=c('#CC0000', '#66CC00', '#0066CC', '#CC6600')
      [unclass(df$CanineGroup)])
```



We can see the relationship between species and the factors using these plots. From `fa.parallel()` function the number of factors recommended are 1 the same number of factors can be verified from the VSS diagram which tells us that there is only 1 factor which helps in identifying the skull measurements. The variables X4, X2, X9 and X6 contribute to factor 1 i.e. RC1

6. Carry out a discriminant function analysis to see how well it is possible to separate the groups using the measurements.

Lets cut the data into two parts 75 % and 25%

```
sample_size <- floor(0.75 * nrow(df))
```

```
train <- sample(nrow(df), size = sample_size)
```

```
train.df <- as.data.frame(df[train, ])
```

```
test.df <- as.data.frame(df[-train, ])
```

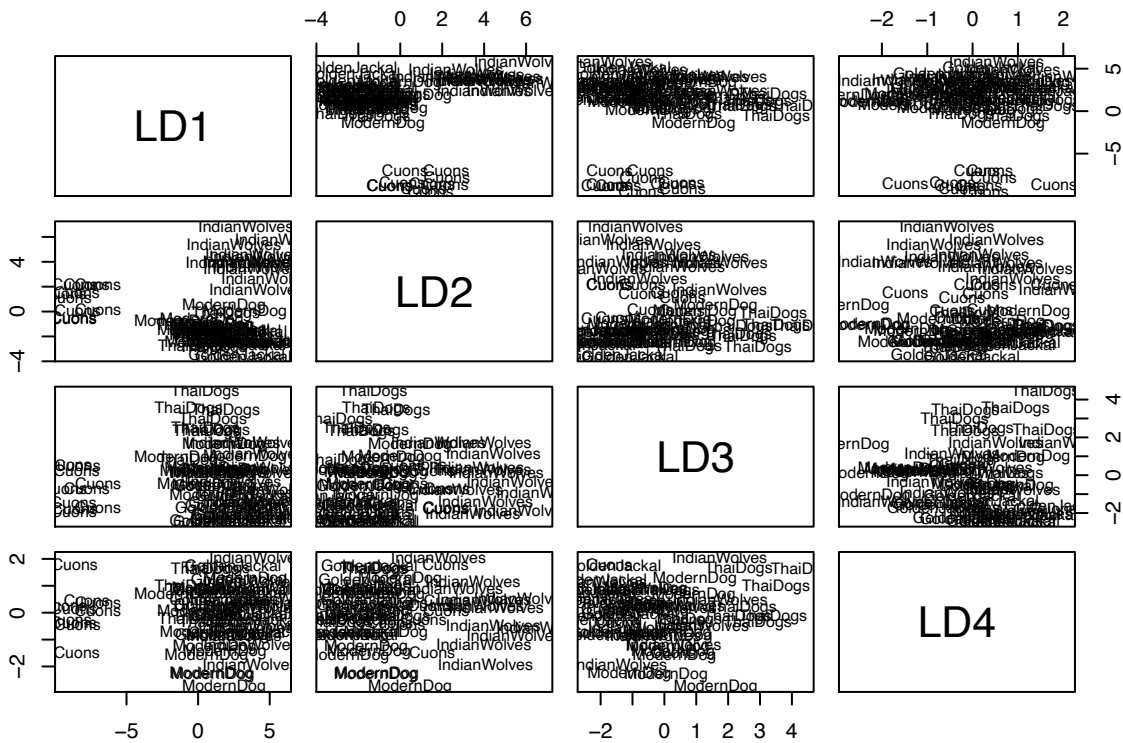
We now have a training and a test set. Training is 75% and test is 25%

```
lda <- lda(formula = train.df$CanineGroup ~  
           X1 + X2 + X3 + X4 + X5 + X6 + X7 + X8 + X9,  
           data = train.df)  
summary(lda)
```

```
##      Length Class  Mode  
## prior      5    -none- numeric  
## counts      5    -none- numeric  
## means     45    -none- numeric  
## scaling    36    -none- numeric  
## lev         5    -none- character  
## svd          4    -none- numeric  
## N            1    -none- numeric  
## call         3    -none- call  
## terms        3    terms  call  
## xlevels      0    -none- list
```

```
plot(lda)
```






```
lda.predict <- predict(lda, newdata = test.df)
```

The above prediction has three items called class, posterior and x

```
lda.data <- cbind(train.df, predict(lda)$x)
```

```
ggplot(lda.data) + geom_point(aes(LD1,
  LD2,
  colour = CanineGroup,
  shape = CanineGroup),
  size = 2.5)
```



```
mean(lda.predict$class == test.df$CanineGroup)
```

```
## [1] 0.85
```

7. investigate each canine group separately to see whether logistic regression shows a significant difference between males and females for the measurements. Note that in view of the small sample sizes available for each group, it is unreasonable to expect to fit a logistic function involving all nine variables with good estimates of parameters. Therefore, consideration should be given to fitting functions using only a subset of the variables.

For Modern Dog

```
xtabs(~ CanineGroup + Gender, df[1:16, ])
```

```
##           Gender
## CanineGroup Female Male Unknown
##   Cuons          0    0      0
## GoldenJackal     0    0      0
## IndianWolves     0    0      0
## ModernDog        8    8      0
## ThaiDogs         0    0      0
```

```
logistic <- glm(Gender ~ X2 + X5 + X7 + X9,
  data = df[1:16, ],
  family = "binomial")
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic)
```

```
##
## Call:
## glm(formula = Gender ~ X2 + X5 + X7 + X9, family = "binomial",
##      data = df[1:16, ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.66110  -0.18145  -0.00007   0.52009   1.55767
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -248.690    197.870  -1.257   0.209
## X2              5.860     5.392   1.087   0.277
## X5             -5.682     5.839  -0.973   0.331
## X7             10.570     8.598   1.229   0.219
## X9             -6.340     5.225  -1.213   0.225
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 22.181  on 15  degrees of freedom
## Residual deviance: 10.377  on 11  degrees of freedom
## AIC: 20.377
##
## Number of Fisher Scoring iterations: 9
```

```
confusion_matrix(logistic)
```

```
##              Predicted Female Predicted Male Total
## Actual Female              7              1      8
## Actual Male              1              7      8
## Total                    8              8     16
```

For Golden Jackals

```
xtabs(~ CanineGroup + Gender, data=df[17:36, ])
```

```
##              Gender
## CanineGroup  Female Male Unknown
## Cuons         0     0     0
## GoldenJackal  10    10     0
## IndianWolves  0     0     0
## ModernDog     0     0     0
## ThaiDogs      0     0     0
```

```
logistic <- glm(Gender ~ X2 + X5 + X7 + X9,
               data = df[17:36, ],
               family = "binomial")
summary(logistic)
```

```
##
## Call:
## glm(formula = Gender ~ X2 + X5 + X7 + X9, family = "binomial",
##      data = df[17:36, ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.35880  -0.26570  -0.00231   0.41771   1.04750
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -94.7912    46.7563  -2.027  0.0426 *
## X2           2.3168     3.3165   0.699  0.4848
## X5           1.3503     1.4360   0.940  0.3470
## X7           0.8786     0.9973   0.881  0.3783
## X9           5.0866     3.9038   1.303  0.1926
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 27.726  on 19  degrees of freedom
## Residual deviance: 10.307  on 15  degrees of freedom
## AIC: 20.307
##
## Number of Fisher Scoring iterations: 6
```

```
confusion_matrix(logistic)
```

```
##              Predicted Female Predicted Male Total
## Actual Female              9              1    10
## Actual Male              0             10    10
## Total                    9             11    20
```

For Cuons

```
xtabs(~ CanineGroup + Gender, data = df[37:53, ])
```

```
##              Gender
## CanineGroup  Female Male Unknown
## Cuons        8      9      0
## GoldenJackal  0      0      0
## IndianWolves  0      0      0
## ModernDog     0      0      0
## ThaiDogs      0      0      0
```

```
logistic <- glm(Gender ~ X2 + X5 + X7 + X9,
                data = df[37:53, ],
                family = "binomial")
summary(logistic)
```

```
##
```

```
## Call:
## glm(formula = Gender ~ X2 + X5 + X7 + X9, family = "binomial",
##      data = df[37:53, ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6647  -1.0491   0.7121   1.0022   1.2719
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -3.3022     12.7101  -0.260   0.795
## X2            -0.6355      1.4351  -0.443   0.658
## X5            -0.9823      1.1581  -0.848   0.396
## X7             0.6569      1.2160   0.540   0.589
## X9             1.8545      2.3458   0.791   0.429
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 23.508  on 16  degrees of freedom
## Residual deviance: 20.665  on 12  degrees of freedom
## AIC: 30.665
##
## Number of Fisher Scoring iterations: 4
```

```
confusion_matrix(logistic)
```

```
##              Predicted Female Predicted Male Total
## Actual Female              4              4      8
## Actual Male              3              6      9
## Total                    7             10     17
```

For Thai Dogs we cannot find and predict the gender since it is unknown in the dataset

For Indian Wolves

```
xtabs(~ CanineGroup + Gender, data = df[64:77, ])
```

```
##              Gender
## CanineGroup  Female Male Unknown
## Cuons        0      0      0
## GoldenJackal 0      0      0
## IndianWolves 6      8      0
## ModernDog     0      0      0
## ThaiDogs      0      0      0
```

```
logistic <- glm(Gender ~ X2 + X5 + X7 + X9,
               data = df[64:77, ],
               family = "binomial")
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logistic)
```

```
##
## Call:
## glm(formula = Gender ~ X2 + X5 + X7 + X9, family = "binomial",
##      data = df[64:77, ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.679e-05 -2.110e-08  2.110e-08  2.110e-08  1.991e-05
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.526e+03  1.212e+07      0      1
## X2           7.368e+01  6.098e+05      0      1
## X5           9.270e+01  1.093e+06      0      1
## X7          -3.977e+01  7.878e+05      0      1
## X9           1.933e+00  1.211e+05      0      1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1.9121e+01  on 13  degrees of freedom
## Residual deviance: 9.2083e-10  on  9  degrees of freedom
## AIC: 10
##
## Number of Fisher Scoring iterations: 25
```

```
confusion_matrix(logistic)
```

```
##              Predicted Female Predicted Male Total
## Actual Female              6              0      6
## Actual Male                0              8      8
## Total                      6              8     14
```

8. Show ROC containing both your discriminant and logistic function for gender classification for the Prehistoric Thai Dog

```
xtabs(~ CanineGroup + Gender, data=df[54:63, ])
```

```
##              Gender
## CanineGroup  Female Male Unknown
## Cuons        0      0      0
## GoldenJackal 0      0      0
## IndianWolves 0      0      0
## ModernDog     0      0      0
## ThaiDogs      0      0     10
```

Randomly assigning the values of Male and Female which is 50% Female and 50% Male

```
df$Gender[54:63] <- c("Female",
                      "Male",
                      "Male",
                      "Female",
                      "Female",
                      "Female",
                      "Male",
                      "Female",
                      "Male",
                      "Male")
```

```
logistic_simple <- glm(Gender ~., data = df[c(-1)], family = "binomial")
summary(logistic_simple)
```

```
##
## Call:
## glm(formula = Gender ~ ., family = "binomial", data = df[c(-1)])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.842  -1.036   0.411   1.050   1.844
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.96963    2.81347  -0.345  0.7304
## X1           0.03666    0.06954   0.527  0.5981
## X2           0.23846    0.49937   0.478  0.6330
## X3           0.10227    0.10040   1.019  0.3084
## X4          -0.17269    0.17928  -0.963  0.3354
## X5          -0.22164    0.27341  -0.811  0.4176
## X6          -0.65196    0.73024  -0.893  0.3720
## X7           0.07159    0.10588   0.676  0.4989
## X8          -0.14762    0.21479  -0.687  0.4919
## X9           1.38744    0.82018   1.692  0.0907 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 106.628  on 76  degrees of freedom
## Residual deviance:  96.677  on 67  degrees of freedom
## AIC: 116.68
##
## Number of Fisher Scoring iterations: 4
```

```
pdata <- predict(logistic_simple, newdata = df, type = "response")
pdata
```

```
##           1           2           3           4           5           6           7           8
## 0.5549875 0.4067477 0.4516714 0.6069229 0.6851925 0.7109144 0.4304144 0.7146370
##           9          10          11          12          13          14          15          16
## 0.5032641 0.5028978 0.7250236 0.4094083 0.5401111 0.5653697 0.5221524 0.5640158
```

```
##      17      18      19      20      21      22      23      24
## 0.518810 0.3248237 0.3205044 0.3938034 0.2959413 0.4753977 0.1825277 0.5729241
##      25      26      27      28      29      30      31      32
## 0.4067892 0.4155722 0.3466740 0.3062581 0.3372643 0.4704012 0.4828238 0.3673306
##      33      34      35      36      37      38      39      40
## 0.3307675 0.3056470 0.3045916 0.2799278 0.5099554 0.7830929 0.7251121 0.6586111
##      41      42      43      44      45      46      47      48
## 0.5764570 0.6187301 0.5476574 0.5842145 0.6186942 0.5397957 0.6831842 0.4922782
##      49      50      51      52      53      54      55      56
## 0.1643165 0.3379583 0.4152510 0.4033627 0.2671415 0.5838307 0.3439986 0.7544320
##      57      58      59      60      61      62      63      64
## 0.4040157 0.5712119 0.3429020 0.5413810 0.5696953 0.6376788 0.4328645 0.6393246
##      65      66      67      68      69      70      71      72
## 0.7712961 0.8331748 0.7224211 0.6925156 0.7347038 0.9189919 0.9728126 0.3885048
##      73      74      75      76      77
## 0.5741460 0.6910565 0.1900079 0.8166125 0.6140271
```

```
pdataF <- as.factor(ifelse(test = as.numeric(pdata > 0.5) == 1,
                           yes = "Female", no = "Male" ))
pdataF
```

```
## [1] Female Male   Male   Female Female Female Male   Female Female Female
## [11] Female Male   Female Female Female Female Female Male   Male   Male
## [21] Male   Male   Male   Female Male   Male   Male   Male   Male   Male
## [31] Male   Male   Male   Male   Male   Male   Female Female Female Female
## [41] Female Female Female Female Female Female Female Male   Male   Male
## [51] Male   Male   Male   Female Male   Female Male   Female Male   Female
## [61] Female Female Male   Female Female Female Female Female Female Female
## [71] Female Male   Female Female Male   Female Female
## Levels: Female Male
```

Here, the actual values will be displayed for the earlier classified Male and Females

```
confusionMatrix(pdataF, as.factor(df$Gender))
```

```
## Warning in levels(reference) != levels(data): longer object length is not a
## multiple of shorter object length

## Warning in confusionMatrix.default(pdataF, as.factor(df$Gender)): Levels are not
## in the same order for reference and data. Refactoring data to match.

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Female Male Unknown
##      Female      16  27      0
##      Male       21  13      0
##      Unknown      0   0      0
##
## Overall Statistics
##
##           Accuracy : 0.3766
```




```
##          95% CI : (0.2687, 0.4944)
##    No Information Rate : 0.5195
##    P-Value [Acc > NIR] : 0.9957
##
##          Kappa : -0.2411
##
##    McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: Female Class: Male Class: Unknown
## Sensitivity          0.4324      0.3250          NA
## Specificity          0.3250      0.4324           1
## Pos Pred Value       0.3721      0.3824          NA
## Neg Pred Value       0.3824      0.3721          NA
## Prevalence           0.4805      0.5195           0
## Detection Rate       0.2078      0.1688           0
## Detection Prevalence 0.5584      0.4416           0
## Balanced Accuracy     0.3787      0.3787          NA
```

```
roc(df$Gender, logistic_simple$fitted.values, plot = TRUE)
```

```
## Setting levels: control = Female, case = Male
```

```
## Setting direction: controls < cases
```




```
##
## Call:
## roc.default(response = df$Gender, predictor = logistic_simple$fitted.values,      plot = TRUE)
##
## Data: logistic_simple$fitted.values in 37 controls (df$Gender Female) < 40 cases (df$Gender Male).
## Area under the curve: 0.7068
```

```
roc(df$Gender,
    logistic_simple$fitted.values,
    plot = TRUE,
    legacy.axes=TRUE,
    xlab="False Positive Percentage",
    ylab="True Postive Percentage",
    col="#377eb8",
    lwd = 4)
```

```
## Setting levels: control = Female, case = Male
```

```
## Setting direction: controls < cases
```



```
##
## Call:
## roc.default(response = df$Gender, predictor = logistic_simple$fitted.values,      plot = TRUE, legacy
##
## Data: logistic_simple$fitted.values in 37 controls (df$Gender Female) < 40 cases (df$Gender Male).
## Area under the curve: 0.7068
```

```
roc.info <- roc(df$Gender, logistic_simple$fitted.values, legacy.axes = TRUE)
```

```
## Setting levels: control = Female, case = Male
```

```
## Setting direction: controls < cases
```

```
str(roc.info)
```

```
## List of 15
## $ percent          : logi FALSE
## $ sensitivities     : num [1:78] 1 1 0.975 0.975 0.975 0.975 0.95 0.95 0.95 0.95 ...
## $ specificities     : num [1:78] 0 0.027 0.027 0.0541 0.0811 ...
## $ thresholds        : num [1:78] -Inf 0.173 0.186 0.229 0.274 ...
## $ direction         : chr "<"
## $ cases             : Named num [1:40] 0.555 0.407 0.452 0.607 0.685 ...
##   .. attr(*, "names")= chr [1:40] "1" "2" "3" "4" ...
## $ controls          : Named num [1:37] 0.503 0.503 0.725 0.409 0.54 ...
##   .. attr(*, "names")= chr [1:37] "9" "10" "11" "12" ...
## $ fun.sesp          :function (thresholds, controls, cases, direction)
## $ auc               : 'auc' num 0.707
##   .. attr(*, "partial.auc")= logi FALSE
##   .. attr(*, "percent")= logi FALSE
##   .. attr(*, "roc")=List of 15
##     ..$ percent          : logi FALSE
##     ..$ sensitivities     : num [1:78] 1 1 0.975 0.975 0.975 0.975 0.95 0.95 0.95 0.95 ...
##     ..$ specificities     : num [1:78] 0 0.027 0.027 0.0541 0.0811 ...
##     ..$ thresholds        : num [1:78] -Inf 0.173 0.186 0.229 0.274 ...
##     ..$ direction         : chr "<"
##     ..$ cases             : Named num [1:40] 0.555 0.407 0.452 0.607 0.685 ...
##     ..   .. attr(*, "names")= chr [1:40] "1" "2" "3" "4" ...
##     ..$ controls          : Named num [1:37] 0.503 0.503 0.725 0.409 0.54 ...
##     ..   .. attr(*, "names")= chr [1:37] "9" "10" "11" "12" ...
##     ..$ fun.sesp          :function (thresholds, controls, cases, direction)
##     ..$ auc               : 'auc' num 0.707
##     ..   .. attr(*, "partial.auc")= logi FALSE
##     ..   .. attr(*, "percent")= logi FALSE
##     ..   .. attr(*, "roc")=List of 8
##       ..$ percent          : logi FALSE
##       ..$ sensitivities     : num [1:78] 1 1 0.975 0.975 0.975 0.975 0.95 0.95 0.95 0.95 ...
##       ..$ specificities     : num [1:78] 0 0.027 0.027 0.0541 0.0811 ...
##       ..$ thresholds        : num [1:78] -Inf 0.173 0.186 0.229 0.274 ...
##       ..$ direction         : chr "<"
##       ..$ cases             : Named num [1:40] 0.555 0.407 0.452 0.607 0.685 ...
##       ..   .. attr(*, "names")= chr [1:40] "1" "2" "3" "4" ...
##       ..$ controls          : Named num [1:37] 0.503 0.503 0.725 0.409 0.54 ...
##       ..   .. attr(*, "names")= chr [1:37] "9" "10" "11" "12" ...
##       ..$ fun.sesp          :function (thresholds, controls, cases, direction)
##       ..   .. attr(*, "class")= chr "roc"
##     ..$ call               : language roc.default(response = df$Gender, predictor = logistic_simple$
##     ..$ original.predictor: Named num [1:77] 0.555 0.407 0.452 0.607 0.685 ...
##     ..   .. attr(*, "names")= chr [1:77] "1" "2" "3" "4" ...
##     ..$ original.response : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 1 1 ...
```

```
## .. ..$ predictor      : Named num [1:77] 0.555 0.407 0.452 0.607 0.685 ...
## .. ..$- attr(*, "names")= chr [1:77] "1" "2" "3" "4" ...
## .. ..$ response       : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 1 1 ...
## .. ..$ levels         : chr [1:2] "Female" "Male"
## .. ..$- attr(*, "class")= chr "roc"
## $ call                : language roc.default(response = df$Gender, predictor = logistic_simple$fitted
## $ original.predictor: Named num [1:77] 0.555 0.407 0.452 0.607 0.685 ...
## ..- attr(*, "names")= chr [1:77] "1" "2" "3" "4" ...
## $ original.response : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 1 1 ...
## $ predictor          : Named num [1:77] 0.555 0.407 0.452 0.607 0.685 ...
## ..- attr(*, "names")= chr [1:77] "1" "2" "3" "4" ...
## $ response           : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 1 1 ...
## $ levels              : chr [1:2] "Female" "Male"
## - attr(*, "class")= chr "roc"
```

```
roc.df <- data.frame(tpp = roc.info$sensitivities * 100,
                     fpp = (1 - roc.info$specificities) * 100,
                     thresholds = roc.info$thresholds)
```

```
head(roc.df)
```

```
##      tpp      fpp thresholds
## 1 100.0 100.00000      -Inf
## 2 100.0  97.29730  0.1734221
## 3  97.5  97.29730  0.1862678
## 4  97.5  94.59459  0.2285747
## 5  97.5  91.89189  0.2735347
## 6  97.5  89.18919  0.2879346
```

```
roc.df[roc.df$tpp > 60 & roc.df$tpp < 80, ]
```

```
##      tpp      fpp thresholds
## 28 77.5 51.35135  0.4229933
## 29 75.0 51.35135  0.4316394
## 30 72.5 51.35135  0.4422679
## 31 70.0 51.35135  0.4610363
## 32 70.0 48.64865  0.4728995
## 33 67.5 48.64865  0.4791108
## 34 67.5 45.94595  0.4875510
## 35 67.5 43.24324  0.4975880
## 36 67.5 40.54054  0.5030810
## 37 67.5 37.83784  0.5066097
## 38 65.0 37.83784  0.5144182
## 39 62.5 37.83784  0.5205167
## 40 62.5 35.13514  0.5309741
## 41 62.5 32.43243  0.5399534
## 42 62.5 29.72973  0.5407460
```

```
roc(df$Gender,
    logistic_simple$fitted.values,
    plot = TRUE,
    legacy.axes = TRUE,
```

```

xlab = "False Positive Percentage",
ylab = "True Postive Percentage",
col = "#377eb8",
lwd = 4)

```

```

## Warning in roc.default(df$Gender, logistic_simple$fitted.values, plot = TRUE, :
## 'response' has more than two levels. Consider setting 'levels' explicitly or
## using 'multiclass.roc' instead

```

```


## Setting levels: control = Female, case = Male

```

```

## Setting direction: controls < cases

```



```

##

```

```

## Call:

```

```

## roc.default(response = df$Gender, predictor = logistic_simple$fitted.values,      plot = TRUE, legacy
##

```

```

## Data: logistic_simple$fitted.values in 37 controls (df$Gender Female) < 40 cases (df$Gender Male).

```

```

## Area under the curve: 0.7068

```

```

par(pty = "s")

```

Conclusion: The Area under the curve is 70% so that means When predicting the gender of thai dogs we are correct 70% of the times

## 9. Predict the Gender for the Prehistoric Thai Dog

a. Explain the reason for choosing the MVA technique for prediction

b. What is the Hit Ratio (Accuracy) of your classification technique?

Let's use Logistic Regression: The gender for the dog will have only two distinct class i.e. either Male or Female, so its very obvious and apt if I use Logistic regression.

```
df_not_thai = df[-c(54:63), -1]
```

```
df_not_thai$Gender = as.factor(df_not_thai$Gender)
```

```
logistic_simple = glm(Gender~., data = df_not_thai, family = "binomial")
summary(logistic_simple)
```

```
##
## Call:
## glm(formula = Gender ~ ., family = "binomial", data = df_not_thai)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9718  -1.1021   0.3837   1.0086   1.7916
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.94892    2.92170  -0.325   0.745
## X1           0.05571    0.07781   0.716   0.474
## X2           0.76742    0.59927   1.281   0.200
## X3           0.09570    0.09692   0.987   0.323
## X4          -0.26780    0.25630  -1.045   0.296
## X5          -0.16317    0.36040  -0.453   0.651
## X6          -0.84409    0.97851  -0.863   0.388
## X7           0.07344    0.13052   0.563   0.574
## X8          -0.24511    0.23483  -1.044   0.297
## X9           1.12319    0.85645   1.311   0.190
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 92.747  on 66  degrees of freedom
## Residual deviance: 81.939  on 57  degrees of freedom
## AIC: 101.94
##
## Number of Fisher Scoring iterations: 4
```

Encoded dependant variable

```
df_not_thai_encoded = df_not_thai
```

```
setDT(df_not_thai_encoded)
```

```
df_not_thai_encoded[Gender == 'Male', 'Gender_flag' := 0]
df_not_thai_encoded[Gender == 'Female', 'Gender_flag' := 1]
str(df_not_thai_encoded)
```

```
## Classes 'data.table' and 'data.frame': 67 obs. of 11 variables:
## $ X1 : num 123 137 121 130 149 125 126 125 121 122 ...
## $ X2 : num 10.1 9.6 10.2 10.7 12 9.5 9.1 9.7 9.6 8.9 ...
## $ X3 : num 23 19 18 24 25 23 20 19 22 20 ...
## $ X4 : num 23 22 21 22 25 20 22 19 20 20 ...
## $ X5 : num 19 19 21 20 21 20 19 19 18 19 ...
## $ X6 : num 7.8 7.8 7.9 7.9 8.4 7.8 7.5 7.5 7.6 7.6 ...
## $ X7 : num 32 32 35 32 35 33 32 32 31 31 ...
## $ X8 : num 33 40 38 37 43 37 35 37 35 35 ...
## $ X9 : num 5.6 5.8 6.2 5.9 6.6 6.3 5.5 6.2 5.3 5.7 ...
## $ Gender : Factor w/ 3 levels "Female","Male",...: 2 2 2 2 2 2 2 2 1 1 ...
## $ Gender_flag: num 0 0 0 0 0 0 0 0 1 1 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
df_not_thai_encoded_logistic = df_not_thai_encoded[, -c('Gender')]

logistic_simple_encoded = glm(Gender_flag~.,
                              data = df_not_thai_encoded_logistic,
                              family = "binomial")
summary(logistic_simple_encoded)
```

```
##
## Call:
## glm(formula = Gender_flag ~ ., family = "binomial", data = df_not_thai_encoded_logistic)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7916  -1.0086  -0.3837   1.1021   1.9718
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.94892    2.92170   0.325   0.745
## X1          -0.05571    0.07781  -0.716   0.474
## X2          -0.76742    0.59927  -1.281   0.200
## X3          -0.09570    0.09692  -0.987   0.323
## X4           0.26780    0.25630   1.045   0.296
## X5           0.16317    0.36040   0.453   0.651
## X6           0.84409    0.97851   0.863   0.388
## X7          -0.07344    0.13052  -0.563   0.574
## X8           0.24511    0.23483   1.044   0.297
## X9          -1.12319    0.85645  -1.311   0.190
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 92.747  on 66  degrees of freedom
## Residual deviance: 81.939  on 57  degrees of freedom
## AIC: 101.94
##
## Number of Fisher Scoring iterations: 4
```

```
confusion_matrix(logistic_simple)
```

```
##           Predicted Female Predicted Male Total
## Actual Female           23           9    32
## Actual Male            12          23    35
## Total                  35          32    67
```

```
confusion_matrix(logistic_simple_encoded)
```


```
##           Predicted 0 Predicted 1 Total
## Actual 0           23          12    35
## Actual 1            9          23    32
## Total              32          35    67
```

```
roc(df_not_thai$Gender,
    logistic_simple$fitted.values,
    plot = TRUE)
```

```
## Warning in roc.default(df_not_thai$Gender, logistic_simple$fitted.values, :
## 'response' has more than two levels. Consider setting 'levels' explicitly or
## using 'multiclass.roc' instead
```

```
## Setting levels: control = Female, case = Male
```

```
## Setting direction: controls < cases
```





```
##
## Call:
## roc.default(response = df_not_thai$Gender, predictor = logistic_simple$fitted.values,      plot = TRUE)
##
## Data: logistic_simple$fitted.values in 32 controls (df_not_thai$Gender Female) < 35 cases (df_not_thai$Gender Male)
## Area under the curve: 0.7214
```


```
par(pty = "s")
```

```
roc(df_not_thai$Gender,
    logistic_simple$fitted.values,
    plot = TRUE)
```

```
## Warning in roc.default(df_not_thai$Gender, logistic_simple$fitted.values, :
## 'response' has more than two levels. Consider setting 'levels' explicitly or
## using 'multiclass.roc' instead
```

```
## Setting levels: control = Female, case = Male
```

```
## Setting direction: controls < cases
```



```
##
## Call:
```

```
## roc.default(response = df_not_thai$Gender, predictor = logistic_simple$fitted.values, plot = TRUE)
##
## Data: logistic_simple$fitted.values in 32 controls (df_not_thai$Gender Female) < 35 cases (df_not_thai$Gender Male)
## Area under the curve: 0.7214
```

Conclusion: From the ROC curve we can conclude with a good AUC of 72.14%

## 10. Create a model to predict length of the Mandible length for Prehistoric Thai Dog.

### a. What is the accuracy of your model

```
ThaiDog<-df[c(54:63), ]
```

X1 is the mandible length

```
head(ThaiDog)
```

```
## # A tibble: 6 x 11
##   CanineGroup   X1    X2    X3    X4    X5    X6    X7    X8    X9 Gender
##   <fct>       <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <fct>
## 1 ThaiDogs    112  10.1   17   18   19   7.7   31   33   5.8 Female
## 2 ThaiDogs    115  10     18   23   20   7.8   33   36   6    Male
## 3 ThaiDogs    136  11.9   22   25   21   8.5   36   39   7    Male
## 4 ThaiDogs    111  9.9    19   20   18   7.3   29   34   5.3 Female
## 5 ThaiDogs    130  11.2   23   27   20   9.1   35   35   6.6 Female
## 6 ThaiDogs    125  10.7   19   26   20   8.4   33   37   6.3 Female
```

```
fit<-lm(X1~., data = ThaiDog[c(-1, -11)])
summary(fit)
```

```
##
## Call:
## lm(formula = X1 ~ ., data = ThaiDog[c(-1, -11)])
##
## Residuals:
##      1      2      3      4      5      6      7      8
## 0.07584 -0.21856 -0.25971 -0.54318 -0.16928 -0.04839  0.08665  1.13774
##      9     10
## -0.41248  0.35137
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -59.9544    34.2225  -1.752   0.330
## X2           -6.9316     3.6587  -1.895   0.309
## X3             4.0239     1.5351   2.621   0.232
## X4            -1.0622     0.6274  -1.693   0.340
## X5            10.1737     5.3594   1.898   0.309
## X6             8.2435     2.7720   2.974   0.207
## X7            -8.2872     3.3399  -2.481   0.244
## X8             1.5138     0.4795   3.157   0.195
```

```
## X9          24.6157      7.8538   3.134    0.197
##
## Residual standard error: 1.429 on 1 degrees of freedom
## Multiple R-squared:  0.9968, Adjusted R-squared:  0.9712
## F-statistic: 38.89 on 8 and 1 DF,  p-value: 0.1234
```

Predicted Values for the variable X1

```
fitted(fit)
```

```
##          1          2          3          4          5          6          7          8
## 111.9242 115.2186 136.2597 111.5432 130.1693 125.0484 131.9133 119.8623
##          9          10
## 122.4125 123.6486
```

```
residuals(fit)
```

```
##          1          2          3          4          5          6
## 0.07584371 -0.21856152 -0.25971481 -0.54317831 -0.16928212 -0.04838838
##          7          8          9          10
## 0.08665477 1.13773883 -0.41248334 0.35137118
```

Anova Table

```
anova(fit) # sum of squares from diff variables
```

```
## Analysis of Variance Table
##
## Response: X1
##      Df Sum Sq Mean Sq F value Pr(>F)
## X2    1 168.189  168.189  82.3414 0.06988 .
## X3    1 113.895  113.895  55.7604 0.08475 .
## X4    1  11.908   11.908   5.8297 0.24997
## X5    1 168.423  168.423  82.4563 0.06983 .
## X6    1  90.833   90.833  44.4697 0.09476 .
## X7    1   9.957    9.957   4.8748 0.27074
## X8    1  52.288   52.288  25.5990 0.12422
## X9    1  20.065   20.065   9.8234 0.19662
## Residuals 1    2.043    2.043
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
gvmodel <- gvlma(fit)
summary(gvmodel)
```

```
##
## Call:
## lm(formula = X1 ~ ., data = ThaiDog[c(-1, -11)])
##
## Residuals:
##          1          2          3          4          5          6          7          8
```

```
## 0.07584 -0.21856 -0.25971 -0.54318 -0.16928 -0.04839 0.08665 1.13774
##          9          10
## -0.41248 0.35137
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -59.9544    34.2225  -1.752   0.330
## X2           -6.9316     3.6587  -1.895   0.309
## X3            4.0239     1.5351   2.621   0.232
## X4           -1.0622     0.6274  -1.693   0.340
## X5           10.1737     5.3594   1.898   0.309
## X6            8.2435     2.7720   2.974   0.207
## X7           -8.2872     3.3399  -2.481   0.244
## X8            1.5138     0.4795   3.157   0.195
## X9           24.6157     7.8538   3.134   0.197
##
## Residual standard error: 1.429 on 1 degrees of freedom
## Multiple R-squared: 0.9968, Adjusted R-squared: 0.9712
## F-statistic: 38.89 on 8 and 1 DF, p-value: 0.1234
##
##
## ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
## USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
## Level of Significance = 0.05
##
## Call:
## gvlma(x = fit)
##
##              Value p-value              Decision
## Global Stat      15.6962 0.003455 Assumptions NOT satisfied!
## Skewness         3.0747 0.079519 Assumptions acceptable.
## Kurtosis         0.7586 0.383763 Assumptions acceptable.
## Link Function    10.0000 0.001565 Assumptions NOT satisfied!
## Heteroscedasticity 1.8628 0.172297 Assumptions acceptable.
```

```
fit1 <- fit
fit2 <- lm(X1 ~ X4 + X6 + X7 + X9, data = ThaiDog)
fit3 <- lm(X1 ~ X2 + X3 + X5 + X6, data = ThaiDog)
```

```
# compare models
anova(fit1, fit2, fit3)
```

```
## Analysis of Variance Table
##
## Model 1: X1 ~ X2 + X3 + X4 + X5 + X6 + X7 + X8 + X9
## Model 2: X1 ~ X4 + X6 + X7 + X9
## Model 3: X1 ~ X2 + X3 + X5 + X6
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      1  2.043
## 2      5 50.777 -4   -48.735 5.9649 0.2968
## 3      5 88.091  0   -37.314
```

```
step <- stepAIC(fit, direction = "both")# step1 take correlation. step2 take highest correlation
```

```
## Start: AIC=2.12
## X1 ~ X2 + X3 + X4 + X5 + X6 + X7 + X8 + X9
##
##      Df Sum of Sq    RSS    AIC
## <none>                2.0426  2.1163
## - X4      1      5.8552  7.8977 13.6399
## - X2      1      7.3315  9.3741 15.3536
## - X5      1      7.3604  9.4030 15.3845
## - X7      1     12.5753 14.6179 19.7966
## - X3      1     14.0345 16.0770 20.7481
## - X6      1     18.0639 20.1064 22.9845
## - X9      1     20.0650 22.1075 23.9333
## - X8      1     20.3561 22.3987 24.0642
```

```
step$anova # display results
```

```
## Stepwise Model Path
## Analysis of Deviance Table
##
## Initial Model:
## X1 ~ X2 + X3 + X4 + X5 + X6 + X7 + X8 + X9
##
## Final Model:
## X1 ~ X2 + X3 + X4 + X5 + X6 + X7 + X8 + X9
##
##
##      Step Df Deviance Resid. Df Resid. Dev    AIC
## 1                1      2.042577  2.11627
```

```
predict.lm(fit1, data.frame(X2 = 10.1,
                             X3 = 17,
                             X4 = 18,
                             X5 = 19,
                             X6 = 7.7,
                             X7 = 31,
                             X8 = 33,
                             X9 = 5.8))
```

```
##      1
## 111.9242
```

```
predict.lm(fit2, data.frame(X2 = 10.1,
                             X3 = 17,
                             X4 = 18,
                             X5 = 19,
                             X6 = 7.7,
                             X7 = 31,
                             X8 = 33,
                             X9 = 5.8))
```

```
##          1
## 115.5767
```

```
predict.lm(fit3, data.frame(X2 = 10.1,
                             X3 = 17,
                             X4 = 18,
                             X5 = 19,
                             X6 = 7.7,
                             X7 = 31,
                             X8 = 33,
                             X9 = 5.8))
```

```
##          1
## 112.4628
```

The predicted value for X1 is the best when predicted with fit 1 we can compare it to the first index since we have tried to predict the first index which had an actual value of 120